



Diffusion Models Presentation

추성재, DGIST


What is Diffusion Model?

- Diffusion Model or Denoising Diffusion Probabilistic Model (DDPM)
- Forward Encoder
 - corrupt training image using multi-step noising process
 - training image turn into sample of Gaussian distribution
- Reverse Decoder
 - deep neural network is trained to invert noising process
 - trained network can generate new images with new sample of Gaussian distribution

What is Diffusion Model?




20.1 Forward Encoder



- Forward process step
 - Markov chain – current state only influenced by previous state
 - step: $\mathbf{z}_t = \sqrt{1 - \beta_t} \mathbf{z}_{t-1} + \sqrt{\beta_t} \epsilon_t$
 where $\epsilon_t \sim \mathcal{N}(\epsilon_t | \mathbf{0}, \mathbf{I})$, $\beta_t \in (0, 1)$, $\beta_1 < \beta_2 < \dots < \beta_T$
- can be rewrite in : $q(\mathbf{z}_t | \mathbf{z}_{t-1}) = \mathcal{N}(\mathbf{z}_t | \sqrt{1 - \beta_t} \mathbf{z}_{t-1}, \beta_t \mathbf{I})$

20.1 Forward Encoder – Diffusion kernel



- Diffusion kernel :

$$q(\mathbf{z}_t | \mathbf{x}) = \mathcal{N}(\mathbf{z}_t | \sqrt{\alpha_t} \mathbf{x}, (1 - \alpha_t) \mathbf{I})$$

where $\alpha_t = \prod_{\tau=1}^t (1 - \beta_\tau)$

can be written like $\mathbf{z}_t = \sqrt{\alpha_t} \mathbf{x} + \sqrt{1 - \alpha_t} \epsilon_t$

if $T \rightarrow \infty$, $q(\mathbf{z}_t | \mathbf{x}) = \mathcal{N}(\mathbf{z}_T | \mathbf{0}, \mathbf{I})$

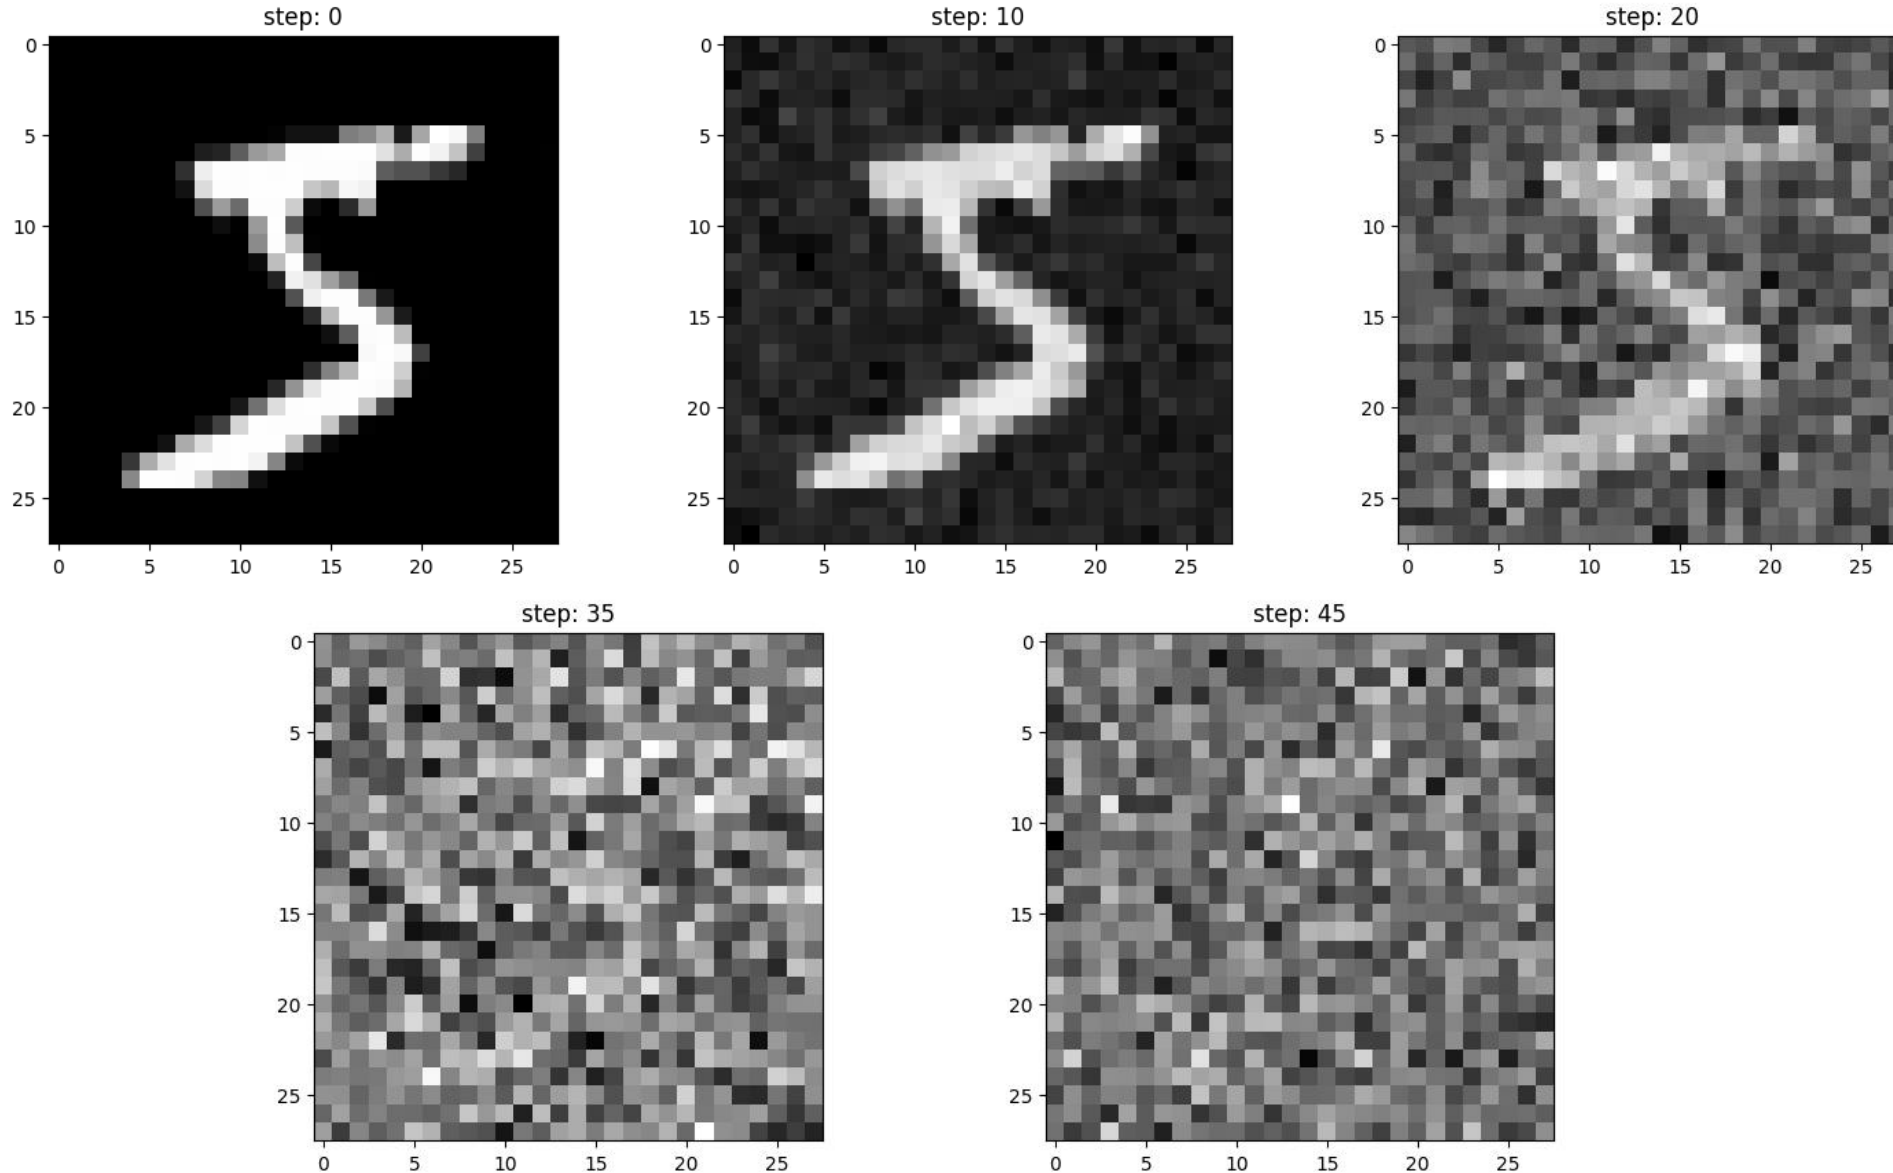
20.1 Forward Encoder – Pytorch example

```
def cal_alpha(beta, step):  
    result = torch.tensor(1).float()  
  
    for i in range(0, step+1):  
        tmp = beta[0, i]  
        result *= (1-tmp)  
  
    return result
```


점진적 forward 과정

```
for t in range(0, terminal_step, 5):  
    noise_t = torch.randn(28,28)  
    alpha_t = cal_alpha(beta, t)  
  
    # forward 과정  
    z_t = forward_image * torch.sqrt(alpha_t) + noise_t * (1-torch.sqrt(alpha_t))  
    plt.figure()  
    plt.imshow(z_t, cmap='gray')
```

20.1 Forward Encoder – Pytorch example



20.1 Forward Encoder – Conditional distribution




$$q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x}) = \frac{q(\mathbf{z}_t|\mathbf{z}_{t-1}, \mathbf{x})q(\mathbf{z}_{t-1}|\mathbf{x})}{q(\mathbf{z}_t|\mathbf{x})}$$

$q(\mathbf{z}_t|\mathbf{z}_{t-1}, \mathbf{x}) = q(\mathbf{z}_t|\mathbf{z}_{t-1})$
 exponential of quadratic form
 -> can be turn into normal distribution

can be ignore
 diffusion kernel: normal distribution


$$q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{z}) = \mathcal{N}(\mathbf{z}_{t-1}|\mathbf{m}(\mathbf{x}, \mathbf{z}_t), \sigma_t^2 \mathbf{I})$$

20.2 Reverse Decoder



- Instead of solving $q(\mathbf{z}_{t-1}|\mathbf{z}_t)$, we learn approximation of reverse distribution: $p(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{w}) = \mathcal{N}(\mathbf{z}_{t-1}|\mu(\mathbf{z}_t, \mathbf{w}, t), \beta_t \mathbf{I})$
- $\mu(\mathbf{z}_t, \mathbf{w}, t)$ is deep neural network with parameter \mathbf{w}
- Main assumption :
 $\beta_t \ll 1$ so that $q(\mathbf{z}_{t-1}|\mathbf{z}_t)$
approximately Gaussian distribution over \mathbf{z}_{t-1}

20.2 Reverse Decoder



- $\mu(\mathbf{z}_t, \mathbf{w}, t)$ can invert total step of forward process
- input dimension and output dimension of $\mu(\mathbf{z}_t, \mathbf{w}, t)$ have to be same, so U-net architecture is commonly used.
- total denoising process:

$$p(\mathbf{x}, \mathbf{z}_1, \dots, \mathbf{z}_T | \mathbf{w}) = p(\mathbf{z}_T) \left\{ \prod_{t=2}^T p(\mathbf{z}_{t-1} | \mathbf{z}_t, \mathbf{w}) \right\} p(\mathbf{x} | \mathbf{z}_1, \mathbf{w}). \quad (20.19)$$

20.2 Reverse Decoder – Evidence Lower Bound (ELBO)

- Total denoising process is intractable: process involve integrating over the neural network functions

- We use Evidence Lower Bound (ELBO)

- for all distribution $q(\mathbf{z})$, following relation always holds

$$\ln p(\mathbf{x}|\mathbf{w}) = \mathcal{L}(\mathbf{w}) + \text{KL}(q(\mathbf{z})||p(\mathbf{z}|\mathbf{x}, \mathbf{w}))$$

- Where Kullback-Leibler divergence term is always bigger than 0, we can get following relation

$$\ln p(\mathbf{x}|\mathbf{w}) \geq \mathcal{L}(\mathbf{w})$$

- Why KL-divergence is always bigger than 0?

$p(\mathbf{z}|\mathbf{x}, \mathbf{w})$ and $q(z)$ is different.

$$\text{so } \text{KL}(q(z)||p(\mathbf{z}|\mathbf{x}, \mathbf{w})) \geq 0$$

- Our objective : train neural network to maximize ELBO

20.2 Reverse Decoder – Evidence Lower Bound (ELBO)

Evidence Lower Bound (ELBO)는 계산이 용이한 임의의 가우시안 분포 $q(z)$ 를 가지고 계산이 불가능한 $\log p_\theta(x)$ (Evidence)의 lower bound를 최대화하는 방식으로 $\log p_\theta(x)$ 을 근사하는 변분 추론 방법입니다.

$$\begin{aligned}\log p_\theta(x) &\geq \sum_z q(z) \log \frac{p_\theta(x, z)}{q(z)} \\ &= \sum_z q(z) \log p_\theta(x, z) - \sum_z q(z) \log q(z) \\ &= \sum_z q(z) \log p_\theta(x, z) + H(q)\end{aligned}$$

$$\therefore ELBO(q) = \sum_z q(z) \log p_\theta(x, z) + H(q)$$

출처 : [생성모델 입문서 - WikiDocs](#)

KL-Divergence

임의의 두 확률 분포 $p(x)$ 와 $q(x)$ 에 대해 KL-Divergence는 다음 성질을 만족합니다.

$$D_{KL}(p||q) = \sum_x p(x) \log \frac{p(x)}{q(x)}$$

- $D_{KL}(p||q) \geq 0$: 두 분포간 차이는 존재해야 합니다 ($\not< 0$)
- $D_{KL}(p||q) = 0$ if $p = q$: 두 분포가 같다면 KLD는 0입니다.
- $D_{KL}(p||q) \neq D_{KL}(q||p)$: 교환법칙은 성립하지 않습니다.

출처 : [생성모델 입문서 - WikiDocs](#)

20.2 Reverse Decoder – Evidence Lower Bound (ELBO)

can be ignored:
has no \mathbf{w}

$$\mathcal{L}(\mathbf{w}) = \mathbb{E}_q \left[\ln \frac{p(\mathbf{z}_T) \left\{ \prod_{t=2}^T p(\mathbf{z}_{t-1} | \mathbf{z}_t, \mathbf{w}) \right\} p(\mathbf{x} | \mathbf{z}_1, \mathbf{w})}{q(\mathbf{z}_1 | \mathbf{x}) \prod_{t=2}^T q(\mathbf{z}_t | \mathbf{z}_{t-1}, \mathbf{x})} \right]$$

$$= \mathbb{E}_q \left[\cancel{\ln p(\mathbf{z}_T)} + \sum_{t=2}^T \ln \frac{p(\mathbf{z}_{t-1} | \mathbf{z}_t, \mathbf{w})}{q(\mathbf{z}_t | \mathbf{z}_{t-1}, \mathbf{x})} - \cancel{\ln q(\mathbf{z}_1 | \mathbf{x})} + \ln p(\mathbf{x} | \mathbf{z}_1, \mathbf{w}) \right] \quad (20.26)$$

can be ignored:
has no \mathbf{w}

$$\mathcal{L}(\mathbf{w}) = \mathbb{E}_q \left[\sum_{t=2}^T \ln \frac{p(\mathbf{z}_{t-1} | \mathbf{z}_t, \mathbf{w})}{q(\mathbf{z}_{t-1} | \mathbf{z}_t, \mathbf{x})} + \ln p(\mathbf{x} | \mathbf{z}_1, \mathbf{w}) \right]$$

$$\mathcal{L}(\mathbf{w}) = \underbrace{\int q(\mathbf{z}_1 | \mathbf{x}) \ln p(\mathbf{x} | \mathbf{z}_1, \mathbf{w}) d\mathbf{z}_1}_{\text{reconstruction term}}$$

$$- \underbrace{\sum_{t=2}^T \int \text{KL}(q(\mathbf{z}_{t-1} | \mathbf{z}_t, \mathbf{x}) || p(\mathbf{z}_{t-1} | \mathbf{z}_t, \mathbf{w})) q(\mathbf{z}_t | \mathbf{x}) d\mathbf{z}_t}_{\text{consistency terms}} \quad (20.32)$$

20.2 Reverse Decoder – Evidence Lower Bound (ELBO)

$$\mathcal{L}(\mathbf{w}) = \underbrace{\int q(\mathbf{z}_1|\mathbf{x}) \ln p(\mathbf{x}|\mathbf{z}_1, \mathbf{w}) d\mathbf{z}_1}_{\text{reconstruction term}} - \sum_{t=2}^T \underbrace{\int \text{KL}(q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x}) \| p(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{w})) q(\mathbf{z}_t|\mathbf{x}) d\mathbf{z}_t}_{\text{consistency terms}} \quad (20.32)$$

$$\begin{aligned} & \text{KL}(q(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{x}) \| p(\mathbf{z}_{t-1}|\mathbf{z}_t, \mathbf{w})) \\ &= \frac{1}{2\beta_t} \|\mathbf{m}_t(\mathbf{x}, \mathbf{z}_t) - \boldsymbol{\mu}(\mathbf{z}_t, \mathbf{w}, t)\|^2 + \text{const} \end{aligned}$$

- KL-divergence term acts like squared-loss function
- minimizing squared error to maximize ELBO

20.2 Reverse Decoder – Predicting the noise

- We introduce new neural network $\mathbf{g}(\mathbf{z}_t, \mathbf{w}, t)$, that **predicts total noise** added to \mathbf{x} to generate \mathbf{z}_t

$$\mu(\mathbf{z}_t, \mathbf{w}, t) = \frac{1}{\sqrt{1 - \beta_t}} \left\{ \mathbf{z}_t - \frac{\beta_t}{\sqrt{1 - \alpha_t}} \mathbf{g}(\mathbf{z}_t, \mathbf{w}, t) \right\}$$

- with several process, we can get

$$\mathcal{L}(\mathbf{w}) = - \sum_{t=1}^T \left\| \mathbf{g}(\sqrt{\alpha_t} \mathbf{x} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t, \mathbf{w}, t) - \boldsymbol{\epsilon}_t \right\|^2$$

Algorithm 20.1: Training a denoising diffusion probabilistic model

Input: Training data $\mathcal{D} = \{\mathbf{x}_n\}$

Noise schedule $\{\beta_1, \dots, \beta_T\}$

Output: Network parameters \mathbf{w}

for $t \in \{1, \dots, T\}$ **do**

$\alpha_t \leftarrow \prod_{\tau=1}^t (1 - \beta_\tau)$ // Calculate alphas from betas

end for

repeat

$\mathbf{x} \sim \mathcal{D}$ // Sample a data point

$t \sim \{1, \dots, T\}$ // Sample a point along the Markov chain

$\boldsymbol{\epsilon} \sim \mathcal{N}(\boldsymbol{\epsilon} | \mathbf{0}, \mathbf{I})$ // Sample a noise vector

$\mathbf{z}_t \leftarrow \sqrt{\alpha_t} \mathbf{x} + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}$ // Evaluate noisy latent variable

$\mathcal{L}(\mathbf{w}) \leftarrow \left\| \mathbf{g}(\mathbf{z}_t, \mathbf{w}, t) - \boldsymbol{\epsilon} \right\|^2$ // Compute loss term

 Take optimizer step

until converged

return \mathbf{w}

20.2 Reverse Decoder – Generating new sample

Algorithm 20.2: Sampling from a denoising diffusion probabilistic model

Input: Trained denoising network $\mathbf{g}(\mathbf{z}, \mathbf{w}, t)$

Noise schedule $\{\beta_1, \dots, \beta_T\}$

Output: Sample vector \mathbf{x} in data space

$\mathbf{z}_T \sim \mathcal{N}(\mathbf{z}|\mathbf{0}, \mathbf{I})$ // Sample from final latent space

for $t \in T, \dots, 2$ **do**

$\alpha_t \leftarrow \prod_{\tau=1}^t (1 - \beta_\tau)$ // Calculate alpha

 // Evaluate network output

$\boldsymbol{\mu}(\mathbf{z}_t, \mathbf{w}, t) \leftarrow \frac{1}{\sqrt{1-\beta_t}} \left\{ \mathbf{z}_t - \frac{\beta_t}{\sqrt{1-\alpha_t}} \mathbf{g}(\mathbf{z}_t, \mathbf{w}, t) \right\}$

$\boldsymbol{\epsilon} \sim \mathcal{N}(\boldsymbol{\epsilon}|\mathbf{0}, \mathbf{I})$ // Sample a noise vector

$\mathbf{z}_{t-1} \leftarrow \boldsymbol{\mu}(\mathbf{z}_t, \mathbf{w}, t) + \sqrt{\beta_t} \boldsymbol{\epsilon}$ // Add scaled noise

end for

$\mathbf{x} = \frac{1}{\sqrt{1-\beta_1}} \left\{ \mathbf{z}_1 - \frac{\beta_1}{\sqrt{1-\alpha_1}} \mathbf{g}(\mathbf{z}_1, \mathbf{w}, t) \right\}$ // Final denoising step

return \mathbf{x}

20.3 Score Matching

- Diffusion model is highly related with “score matching” deep generative model
- “score matching” model uses “score function” or “Stein score”

$$\mathbf{s}(\mathbf{x}) = \nabla_{\mathbf{x}} \ln p(\mathbf{x})$$

- $\mathbf{s}(\mathbf{x})$ has same dimensionality as \mathbf{x}

20.3 Score Matching – score loss function

- loss function to learn $s(\mathbf{x}, \mathbf{w})$

$$J(\mathbf{w}) = \frac{1}{2} \int \|\mathbf{s}(\mathbf{x}, \mathbf{w}) - \nabla_{\mathbf{x}} \ln p(\mathbf{x})\|^2 p(\mathbf{x}) d\mathbf{x}$$

- Problem of this loss function is we do not know $\nabla_{\mathbf{x}} \ln p(\mathbf{x})$
- So, we need to use data points “smear out” from the noise kernel (Parzen estimator)

$$q_{\sigma}(\mathbf{z}) = \int q(\mathbf{z}|\mathbf{x}, \sigma) p(\mathbf{x}) d\mathbf{x} \quad (20.47)$$

where $q(\mathbf{z}|\mathbf{x}, \sigma)$ is the *noise kernel*. A common choice of kernel is the Gaussian

$$q(\mathbf{z}|\mathbf{x}, \sigma) = \mathcal{N}(\mathbf{z}|\mathbf{x}, \sigma^2 \mathbf{I}). \quad (20.48)$$

20.3 Score Matching – score loss function

- Change the loss function:

$$J(\mathbf{w}) = \frac{1}{2} \iint \|\mathbf{s}(\mathbf{z}, \mathbf{w}) - \nabla_{\mathbf{z}} \ln q(\mathbf{z}|\mathbf{x}, \sigma)\|^2 q(\mathbf{z}|\mathbf{x}, \sigma) p(\mathbf{x}) d\mathbf{z} d\mathbf{x} + \text{const.}$$

- If we choose noise kernel as Gaussian kernel, $\nabla_{\mathbf{z}} \ln q(\mathbf{z}|\mathbf{x}, \sigma) = -\frac{1}{\sigma} \epsilon$
- Then, loss function measures **difference between neural network and noise**
- Therefore, this loss function has same minimum as **diffusion model's loss function** : $\mathbf{s}(\mathbf{x}, \mathbf{w}) \sim -\frac{1}{\sigma} \mathbf{g}(\mathbf{z}, \mathbf{w})$
- Langevin dynamics is used to sample from score function

20.3 Score Matching – Noise variance

- Score function has three problems
 - If probability density is zero, score function is undefined
 - If data density is low, estimation of score function is inaccurate
 - If data distribution is mixture of disjoint distribution, Langevin procedure might not sample properly
- All problems can be solved with choosing large noise variance σ^2 , but this might cause distortion of original distribution
- This can be handled with sequence of noise variance which follows

$$\sigma_1^2 < \sigma_2^2 < \dots < \sigma_T^2$$

- score network is modified to get noise variance

$$s(\mathbf{x}, \mathbf{w}, \sigma^2)$$

20.4 Guided Diffusion

- In many applications, we want diffusion models to generate image with certain “conditions”
 - conditions could be label or text description (prompt)
- most simple approach is to add conditioning variable \mathbf{c} inside denoising neural network $g(\mathbf{z}, \mathbf{w}, t, \mathbf{c})$ and train network with $\{\mathbf{x}_n, \mathbf{c}_n\}$
- In this approach, we need “guidance” which is weight given to conditioning variable

20.4 Guided Diffusion – Classifier guidance

- If there are trained classifier $p(\mathbf{c}|\mathbf{x})$, and there are diffusion model using score function, there are relation that

$$\begin{aligned}
 \nabla_{\mathbf{x}} \ln p(\mathbf{x}|\mathbf{c}) &= \nabla_{\mathbf{x}} \ln \left\{ \frac{p(\mathbf{c}|\mathbf{x})p(\mathbf{x})}{\cancel{p(\mathbf{c})}} \right\} \xrightarrow{\text{diffusion model using score function}} \nabla_{\mathbf{x}} \ln p(\mathbf{c}) = 0 \\
 &= \nabla_{\mathbf{x}} \ln p(\mathbf{x}) + \nabla_{\mathbf{x}} \ln p(\mathbf{c}|\mathbf{x}) \xrightarrow{\text{makes denoising process maximize probability of } \mathbf{c}}
 \end{aligned}$$

- add hyperparameter λ (guidance scale), score function becomes

$$\text{score}(\mathbf{x}, \mathbf{c}, \lambda) = \nabla_{\mathbf{x}} \ln p(\mathbf{x}) + \lambda \nabla_{\mathbf{x}} \ln p(\mathbf{c}|\mathbf{x}).$$

- drawback: need to train separate classifier $p(\mathbf{c}|\mathbf{x})$

20.4 Guided Diffusion – Classifier-free guidance

- If we change score function of “Classifier guidance”

$$\text{score}(\mathbf{x}, \mathbf{c}, \lambda) = \lambda \nabla_{\mathbf{x}} \ln p(\mathbf{x}|\mathbf{c}) + (1 - \lambda) \nabla_{\mathbf{x}} \ln p(\mathbf{x})$$

- We can avoid training separate networks. ($p(\mathbf{x}) = p(\mathbf{x}|\mathbf{c} = \mathbf{0})$)
- This approach is used in many guided diffusion models
 - Text-guided diffusion model : using prompt (text sequence) as conditioning variable
 - image super-resolution : denoising high-resolution image with using low-resolution image as conditioning variable

20.4 Guided Diffusion – Stable Diffusion

